

# Burning Problem using ML

Imitation Learning for GNNs


December 4, 2025

# Introduction


# The Burning Problem

- ▶ **Graph Burning Problem:** equivalent to the **Graph Covering Problem**:
  - ▶ A ball of radius  $r$  around a node covers all nodes of distance  $\leq r$ .
  - ▶ The goal is to determine the shortest sequence of balls to cover the entire graph.
- ▶ **Burning number** = minimum number of balls used to burn the graph.


## Step 0:




# Step 1:




## Step 2:



## Step 3:



## Step 4:



# Motivation

- ▶ **NP-hard problem:** Finding the burning number is computationally intractable for large graphs.
- ▶ Optimal solutions are hard to find, so we often rely on heuristics – it would be nice to have a way to reliably find good coverings.
- ▶ **Real-world relevance:** models spread of information, diseases, or influence in networks.

## Problem Formulation

- ▶ **Problem:** Can we train a model to efficiently find a minimal solution to the **Burning Problem**?
- ▶ **Inspiration:** Prior work showed success on another NP-hard task (finding Hamiltonian cycles using graph neural networks).
- ▶ **Objective:** We aim to apply a similar approach (imitation learning + GNN) to the graph burning problem.

## Problem Simplification

- ▶ We simplified the problem to covering trees.
- ▶ Trees are a practical choice as they are easy to generate uniformly and they are the hardest to cover in practice (not a well connected network).
- ▶ Furthermore, they have good generalization property since to cover a graph we are only required to cover one of its spanning trees.

# Data Generation


## Data Generation

- ▶ We decided to generate our own dataset to have more control over its properties and ensure that the training set is well representative of general network structures.
- ▶ **Uniform random trees** generated via *Prüfer sequences*: each random Prüfer code of length  $n - 2$  yields a uniformly random labeled tree. These sequences can be easily generated and decoded back to their tree representation.

# Methodology

- ▶ For each tree, computed the exact **burning number** by brute force:
  - ▶ Tried all possible graph covers and chose a minimal one arbitrarily.
  - ▶ Exponential search – very slow: computing optimal burns for 5,000 trees took over 10 hours.
- ▶ These optimal sequences (node orderings) serve as expert demonstrations for training.

# Examples



# Our Model

## Imitation Learning Approach

- ▶ We adopt **imitation learning**: train a policy to mimic an expert (the minimal covering).
- ▶ Formulate as a sequential decision process:
  - ▶ State = graph + inherit graph property  $h$  + partial solution  $\mathbf{x}$
  - ▶ Action = estimate the probabilities  $p$  of each node being the next center and pick the node  $c$  with largest associated probability.
  - ▶ Expert policy: compare with the cover label and punish the model by  $-\ln(p[c])$ . We reinitialize  $\mathbf{x}$  with the correct partial solution and update  $h$  with model output.

## GNN and message passing layers

- ▶ Nodes exchange information with neighbors through message passing layers, aggregating structural context.
- ▶ After  $L$  layers, each node's embedding captures information from its  $L$ -neighborhood (graph structure + covers).
- ▶ This is the entire essence of the model and how we may hope to correctly choose optimal centers.

# Algorithm

## Input:

$$\begin{cases} G & \text{graph with } n \text{ vertices} \\ x \in \mathbb{R}^{n \times d_{\text{in}}} & \text{partial solution representation} \\ h \in \mathbb{R}^{n \times d_h} & \text{latent features} \end{cases}$$

## Output:

$$p \in [0, 1]^n \quad \text{next-step probabilities per node}$$

## Hyperparameters:

$$d_{\text{in}} = 2, \quad d_h = 32, \quad n_p = 4$$

## Parameters:

$$\theta \equiv \{W_E, b_E, W_P, b_P, \dots\} \quad \text{neural network weights}$$

# Forwarding

## Algorithm:

*Encoder - Initialize features:*

$$z_i = W_E(x_i \oplus h_i) + b_E \in \mathbb{R}^{d_h}, \quad h_i = z_i$$

*Message-passing - Apply residual max-MPNN layers:*

for  $k = 1$  to  $n_p$  :

$$\begin{cases} m_i = \max_{j \sim i} \text{ReLU} \left( W_M^{(k)}(h_i \oplus h_j) + b_M \right) \in \mathbb{R}^{d_h} \\ h_i = h_i + \text{ReLU} \left( W^{(k)}(h_i \oplus m_i) + b^{(k)} \right) \in \mathbb{R}^{d_h} \end{cases}$$

*Decoder - Extract logits and probabilities:*

$$l_i = W_D(z_i \oplus h_i) + b_D \in \mathbb{R}, \quad i = 1, \dots, n$$

if  $i$  is already a center,  $l_i = -\infty$

$$p = \text{softmax}(l) \in \mathbb{R}^n$$

# GNN Architecture




Figure: MP layer 1-node

# Training

## Training and Experiments

- ▶ Training via supervised learning (behavior cloning): minimize error between predicted next-node and expert's choice.
- ▶ We used Adam optimizer with steps of size  $10^{-3}$  for 100 epochs.

# Training Algorithm

```
1: for epoch = 1 to max do
2:   for tree,cover  $\in D$  do
3:     h = Initializeh(tree)
4:     loss = 0
5:     for k = 1 to len(cover) do
6:       x = EncodeCoverk(cover,k)
7:       c = center of k-th ball
8:       p,h = model(tree,x,h)
9:       loss += -ln(p[c])
10:      end for
11:      backprop(loss)
12:    end for
13:  end for
```


# Improving Training

- ▶ An important aspect of the model is the initialization of  $h$ . We tested multiple different approaches which include adding different structural graph properties, initializing to a zero matrix and adding random noise to the matrix.
- ▶ The optimal and most intuitive result was to initialize  $h$  to be a column of the node degrees and eigenvectors of the Laplacian while completing to  $n \times 32$  matrix with zeros.

# Improving Training

- ▶ We also experimented on the number of message passing layers, we concluded that 4 was optimal. This is also intuitively correct since the balls we use in the coverings are (conjectured to be) at most of size 4.
- ▶ Pool function choice, sum, average and max.

# Loss Over Time



# Results

## Evaluation Procedure

- ▶ After training, we evaluate the model on 1,000 unseen trees. For each test tree, we retrieve the optimal burning number  $k$  computed by brute-force.
- ▶ We run the model iteratively  $k$  times, updating the input at each step by adding the ball chosen by the model previously.
- ▶ After  $k$  iterations, we verify if the balls placed cover the entire graph, indicating a valid solution.

## Performance and Success Criteria

- ▶ The model is **not** evaluated by reproducing the exact brute-force solution.
- ▶ Success means the solution covers the graph using exactly  $k$  balls.
- ▶ Our model achieves a success rate of **22%**.
- ▶ Randomly sampling  $k$  min centers yields only about **0.3%** success.

Method	Success Rate
Random Sampling	0.3%
Our Model	22%


Table: Success rate comparison

## Interpretation and Generalization


- ▶ Solutions from our model and brute-force are rarely identical but show strong similarities.
- ▶ This indicates the model has **actually** learned the underlying problem structure.
- ▶ Rather than *simply* imitating expert solutions, the model generalizes well.

# Brute Force vs Model solutions (1)

Model Solution (Graph 1)




Algorithm Solution (Graph 1)




# Brute Force vs Model solutions (2)



Model Solution (Graph 33)



Algorithm Solution (Graph 33)




# Brute Force vs Model solutions (3)




# Brute Force vs Model solutions (4)

Model Solution (Graph 53)




Algorithm Solution (Graph 53)




# Brute Force vs Model solutions (5)

Model Solution (Graph 7)



Algorithm Solution (Graph 7)



# Conclusion

# Conclusion & Key Results

**Goal:** Solve the NP-hard **Graph Burning Problem** using GNNs and imitation learning.

► **Approach:**

- Trained a GNN to mimic brute-force solutions on **uniform random trees**.
- Used message-passing to iteratively select burning centers.

► **Results:**

- Model achieved **22% success rate** (vs. **0.3%** for random sampling).
- Learned to generalize: solutions were structurally similar to brute-force (though not identical).

## Discussion 1: Performance Gap with Reference Model

- ▶ Reference model finds Hamiltonian cycles in **75% of cases** for graphs of size 25, vs **22%** for us.
- ▶ Our model underperforms due to key differences:
  - ▶ **Nature of the problem:**
    - ▶ Hamiltonian cycle task has consistent output at each step.
    - ▶ Our task: ball radius increases with each step  $\Rightarrow$  changing output.
  - ▶ **Training becomes harder** due to the shifting objective.
  - ▶ **Fewer model calls per graph:**
    - ▶ Hamiltonian: 25 calls/graph (one per node).
    - ▶ Our task: typically  $\leq 4$  calls/graph.
  - ▶ **Less opportunity to learn latent structure.**

## Discussion 2: Next Steps and Generalization

- ▶ Inspired by the Hamiltonian model's scalability to larger graphs, we could evaluate our model's generalization ability.
- ▶ Proposed experiment:
  - ▶ Generate a dataset of 1,000 trees, each with 50 nodes.
  - ▶ Use brute-force to compute **ground truth** solutions.
  - ▶ Measure model's success rate in matching optimal outputs.
- ▶ **High success rate** would:
  - ▶ Indicate strong learning and generalization.
  - ▶ Validate our model beyond its training data.
- Was not possible for us given the enormous time necessary to compute brute-force solutions.

# References

-  Bonato, A., Janssen, J., & Roshanbin, E. (2014). *Burning a Graph as a Model of Social Contagion*. Algorithms and Models for the Web Graph (WAW 2014), LNCS 8882, pp. 13–22.
-  Bosnić, F., & Šikić, M. (2023). *Finding Hamiltonian cycles with graph neural networks*. arXiv:2306.06523.